

Week 4 - Regression Analysis - Python-HR_Data

October 3, 2018

1 Data Warehousing and Data Mining

1.1 Labs

1.1.1 Prepared by Gilroy Gordon

Contact Information SCIT ext. 3643

ggordon@utech@gmail.com

gilroy.gordon@utech.edu.jm

1.1.2 Week 3 - Regression Analysis in Python

Additional Reference Resources:

http://scikit-learn.org/stable/modules/linear_model.html

1.2 Objectives

- > Data Transformation
- > Data Mining
 - > Linear Regression
- > Model Evaluation and Prediction
 - > Train/Test Split - 70/30
- > Presentation
 - > Scatter Plot

1.3 Import required libraries and acquire data

NB. The data required was retrieved from the required text for this course. This should assist you in following the concepts from the book better

```
In [1]: # import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: data_path = './data/hr_data.csv' # Path to data file
data = pd.read_csv(data_path)
data.head(15)
```

```
Out[2]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	\
0	0.38	0.53	2	157	
1	0.80	0.86	5	262	
2	0.11	0.88	7	272	
3	0.72	0.87	5	223	
4	0.37	0.52	2	159	
5	0.41	0.50	2	153	
6	0.10	0.77	6	247	
7	0.92	0.85	5	259	
8	0.89	1.00	5	224	
9	0.42	0.53	2	142	
10	0.45	0.54	2	135	
11	0.11	0.81	6	305	
12	0.84	0.92	4	234	
13	0.41	0.55	2	148	
14	0.36	0.56	2	137	

	time_spend_company	Work_accident	left	promotion_last_5years	sales	\
0	3	0	1	0	sales	
1	6	0	1	0	sales	
2	4	0	1	0	sales	
3	5	0	1	0	sales	
4	3	0	1	0	sales	
5	3	0	1	0	sales	
6	4	0	1	0	sales	
7	5	0	1	0	sales	
8	5	0	1	0	sales	
9	3	0	1	0	sales	
10	3	0	1	0	sales	
11	4	0	1	0	sales	
12	5	0	1	0	sales	
13	3	0	1	0	sales	
14	3	0	1	0	sales	

	salary
0	low
1	medium
2	medium
3	low
4	low
5	low
6	low
7	low
8	low

```

9      low
10     low
11     low
12     low
13     low
14     low

```

```
In [3]: # What columns are in the data set ? Do they have spaces that I should consider
data.columns
```

```
Out[3]: Index(['satisfaction_level', 'last_evaluation', 'number_project',
              'average_monthly_hours', 'time_spend_company', 'Work_accident', 'left',
              'promotion_last_5years', 'sales', 'salary'],
              dtype='object')
```

1.4 Aim: Can we determine a person's Satisfaction Level based on the other factors?

age = a(last_evaluation) + b(number_project) + c(average_monthly_hours) + d(time_spend_company)

The coefficients a-d, what are they? What is the relationship between the variables? Does multicollinearity exist?

I have created a function below `create_label_encoder_dict` to assist with this. The function accepts a dataframe object and uses the `LabelEncoder` class from `sklearn.preprocessing` to encode (dummy encoding) or transform non-numerical columns to numbers. Finally it returns a dictionary object of all the encoders created for each column.

The `LabelEncoder` is a useful resource as it not only automatically transforms all values in a column but also keeps a track of what values were transformed from. i.e. It will change all Female to 0 and all Male to 1

```
In [4]: def create_label_encoder_dict(df):
        from sklearn.preprocessing import LabelEncoder

        label_encoder_dict = {}
        for column in df.columns:
            # Only create encoder for categorical data types
            if not np.issubdtype(df[column].dtype, np.number) and column != 'Age':
                label_encoder_dict[column] = LabelEncoder().fit(df[column])
        return label_encoder_dict

In [5]: label_encoders = create_label_encoder_dict(data)
        print("Encoded Values for each Label")
        print("="*32)
        for column in label_encoders:
            print("="*32)
            print('Encoder(%s) = %s' % (column, label_encoders[column].classes_ ))
            print(pd.DataFrame([range(0, len(label_encoders[column].classes_))], columns=label_en
```

Encoded Values for each Label

=====

=====

```
Encoder(sales) = ['IT' 'RandD' 'accounting' 'hr' 'management' 'marketing' 'product_mng'
                  'sales' 'support' 'technical']
```

Encoded Values

IT	0
RandD	1
accounting	2
hr	3
management	4
marketing	5
product_mng	6
sales	7
support	8
technical	9

=====

```
Encoder(salary) = ['high' 'low' 'medium']
```

Encoded Values

high	0
low	1
medium	2

```
In [6]: # Apply each encoder to the data set to obtain transformed values
data2 = data.copy() # create copy of initial data set
for column in data2.columns:
    if column in label_encoders:
        data2[column] = label_encoders[column].transform(data2[column])

print("Transformed data set")
print("="*32)
data2
```

Transformed data set

=====

```
Out[6]:
```

	satisfaction_level	last_evaluation	number_project	\
0	0.38	0.53	2	
1	0.80	0.86	5	
2	0.11	0.88	7	
3	0.72	0.87	5	
4	0.37	0.52	2	
5	0.41	0.50	2	
6	0.10	0.77	6	
7	0.92	0.85	5	
8	0.89	1.00	5	

9	0.42	0.53	2
10	0.45	0.54	2
11	0.11	0.81	6
12	0.84	0.92	4
13	0.41	0.55	2
14	0.36	0.56	2
15	0.38	0.54	2
16	0.45	0.47	2
17	0.78	0.99	4
18	0.45	0.51	2
19	0.76	0.89	5
20	0.11	0.83	6
21	0.38	0.55	2
22	0.09	0.95	6
23	0.46	0.57	2
24	0.40	0.53	2
25	0.89	0.92	5
26	0.82	0.87	4
27	0.40	0.49	2
28	0.41	0.46	2
29	0.38	0.50	2
...
14969	0.43	0.46	2
14970	0.78	0.93	4
14971	0.39	0.45	2
14972	0.11	0.97	6
14973	0.36	0.52	2
14974	0.36	0.54	2
14975	0.10	0.79	7
14976	0.40	0.47	2
14977	0.81	0.85	4
14978	0.40	0.47	2
14979	0.09	0.93	6
14980	0.76	0.89	5
14981	0.73	0.93	5
14982	0.38	0.49	2
14983	0.72	0.84	5
14984	0.40	0.56	2
14985	0.91	0.99	5
14986	0.85	0.85	4
14987	0.90	0.70	5
14988	0.46	0.55	2
14989	0.43	0.57	2
14990	0.89	0.88	5
14991	0.09	0.81	6
14992	0.40	0.48	2
14993	0.76	0.83	6
14994	0.40	0.57	2

14995	0.37	0.48	2
14996	0.37	0.53	2
14997	0.11	0.96	6
14998	0.37	0.52	2

	average_monthly_hours	time_spend_company	Work_accident	left	\
0	157	3	0	1	
1	262	6	0	1	
2	272	4	0	1	
3	223	5	0	1	
4	159	3	0	1	
5	153	3	0	1	
6	247	4	0	1	
7	259	5	0	1	
8	224	5	0	1	
9	142	3	0	1	
10	135	3	0	1	
11	305	4	0	1	
12	234	5	0	1	
13	148	3	0	1	
14	137	3	0	1	
15	143	3	0	1	
16	160	3	0	1	
17	255	6	0	1	
18	160	3	1	1	
19	262	5	0	1	
20	282	4	0	1	
21	147	3	0	1	
22	304	4	0	1	
23	139	3	0	1	
24	158	3	0	1	
25	242	5	0	1	
26	239	5	0	1	
27	135	3	0	1	
28	128	3	0	1	
29	132	3	0	1	
...	
14969	157	3	0	1	
14970	225	5	0	1	
14971	140	3	0	1	
14972	310	4	0	1	
14973	143	3	0	1	
14974	153	3	0	1	
14975	310	4	0	1	
14976	136	3	0	1	
14977	251	6	0	1	
14978	144	3	0	1	
14979	296	4	0	1	

14980	238	5	0	1
14981	162	4	0	1
14982	137	3	0	1
14983	257	5	0	1
14984	148	3	0	1
14985	254	5	0	1
14986	247	6	0	1
14987	206	4	0	1
14988	145	3	0	1
14989	159	3	1	1
14990	228	5	1	1
14991	257	4	0	1
14992	155	3	0	1
14993	293	6	0	1
14994	151	3	0	1
14995	160	3	0	1
14996	143	3	0	1
14997	280	4	0	1
14998	158	3	0	1

	promotion_last_5years	sales	salary
0	0	7	1
1	0	7	2
2	0	7	2
3	0	7	1
4	0	7	1
5	0	7	1
6	0	7	1
7	0	7	1
8	0	7	1
9	0	7	1
10	0	7	1
11	0	7	1
12	0	7	1
13	0	7	1
14	0	7	1
15	0	7	1
16	0	7	1
17	0	7	1
18	1	7	1
19	0	7	1
20	0	7	1
21	0	7	1
22	0	7	1
23	0	7	1
24	0	7	1
25	0	7	1
26	0	7	1

27	0	7	1
28	0	2	1
29	0	2	1
...
14969	0	7	2
14970	0	7	2
14971	0	7	2
14972	0	2	2
14973	0	2	2
14974	0	2	2
14975	0	3	2
14976	0	3	2
14977	0	3	2
14978	0	3	2
14979	0	9	2
14980	0	9	0
14981	0	9	1
14982	0	9	2
14983	0	9	2
14984	0	9	2
14985	0	9	2
14986	0	9	1
14987	0	9	1
14988	0	9	1
14989	0	9	1
14990	0	8	1
14991	0	8	1
14992	0	8	1
14993	0	8	1
14994	0	8	1
14995	0	8	1
14996	0	8	1
14997	0	8	1
14998	0	8	1

[14999 rows x 10 columns]

```
In [7]: # separate our data into dependent (Y) and independent(X) variables
X_data = data2[['last_evaluation', 'number_project', 'average_monthly_hours', 'time_spent_company']]
Y_data = data2['satisfaction_level']
```

1.5 70/30 Train Test Split

We will split the data using a 70/30 split. i.e. 70% of the data will be randomly chosen to train the model and 30% will be used to evaluate the model

```
In [8]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_data, Y_data, test_size=0.30)
```



```

In [9]: # Import linear model package (has several regression classes)
        from sklearn import linear_model

In [10]: # Create an instance of linear regression
         reg = linear_model.LinearRegression()

In [11]: reg.fit(X_train,y_train)

Out[11]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)

In [12]: reg.coef_

Out[12]: array([ 2.52091853e-01, -3.86756827e-02,  1.02855525e-04, -1.40672606e-02])

In [13]: X_train.columns

Out[13]: Index(['last_evaluation', 'number_project', 'average_monthly_hours',
               'time_spend_company'],
               dtype='object')

In [14]: print("Regression Coefficients")
         pd.DataFrame(reg.coef_,index=X_train.columns,columns=["Coefficient"])

Regression Coefficients

Out[14]:
               Coefficient
last_evaluation      0.252092
number_project     -0.038676
average_monthly_hours  0.000103
time_spend_company  -0.014067

In [15]: # Intercept
         reg.intercept_

Out[15]: 0.6078482685006579

In [16]: # Make predictions using the testing set
         test_predicted = reg.predict(X_test)
         test_predicted

Out[16]: array([0.60151576, 0.64939671, 0.52336411, ..., 0.67138193, 0.69110759,
               0.5978772 ])

In [17]: data3 = X_test.copy()
         data3['predicted_satisfaction_level']=test_predicted
         data3['satisfaction_level']=y_test
         data3.head()

```

```
Out[17]:
```

	last_evaluation	number_project	average_monthly_hours	\
14667	0.93	5	223	
13373	0.89	4	137	
9246	0.55	5	121	
12261	0.97	5	263	
9359	0.54	3	271	

	time_spend_company	predicted_satisfaction_level	satisfaction_level
14667	5	0.601516	0.86
13373	3	0.649397	0.78
9246	3	0.523364	0.27
12261	5	0.615714	0.82
9359	3	0.613623	0.97

```
In [18]: from sklearn.metrics import mean_squared_error, r2_score
```

```
In [19]: # The mean squared error
print("Mean squared error: %.2f" % mean_squared_error(y_test, test_predicted))
```

Mean squared error: 0.06

```
In [20]: # Explained variance score: 1 is perfect prediction
# R squared
print('Variance score: %.2f' % r2_score(y_test, test_predicted))
```

Variance score: 0.06

```
In [21]: help(reg.score)
```

Help on method score in module sklearn.base:

score(X, y, sample_weight=None) method of sklearn.linear_model.base.LinearRegression instance
Returns the coefficient of determination R^2 of the prediction.

The coefficient R^2 is defined as $(1 - u/v)$, where u is the residual sum of squares $((y_{\text{true}} - y_{\text{pred}}) ** 2).sum()$ and v is the total sum of squares $((y_{\text{true}} - y_{\text{true}.mean()}) ** 2).sum()$.

The best possible score is 1.0 and it can be negative (because the model can be arbitrarily worse). A constant model that always predicts the expected value of y , disregarding the input features, would get a R^2 score of 0.0.

Parameters

X : array-like, shape = (n_samples, n_features)
Test samples.

```
y : array-like, shape = (n_samples) or (n_samples, n_outputs)
    True values for X.
```

```
sample_weight : array-like, shape = [n_samples], optional
    Sample weights.
```

Returns

```
score : float
    R2 of self.predict(X) wrt. y.
```

```
In [22]: reg.score(X_test,y_test)
```

```
Out[22]: 0.061684941263884914
```

1.5.1 Visualizations

It's difficult to plot a scatter plot with so many dimensions

How about Dimensionality Reduction?

One such method - Principal Component Analysis

```
In [23]: from sklearn.decomposition import PCA
```

```
In [24]: pca = PCA(n_components=1)
```

```
In [25]: pca.fit(data2[X_train.columns])
```

```
Out[25]: PCA(copy=True, iterated_power='auto', n_components=1, random_state=None,
    svd_solver='auto', tol=0.0, whiten=False)
```

```
In [26]: pca.components_
```

```
Out[26]: array([[0.00116455, 0.01030169, 0.99993927, 0.00373904]])
```

```
In [27]: pca.n_features_
```

```
Out[27]: 4
```

```
In [28]: pca.n_components_
```

```
Out[28]: 1
```

Now that we can reduce our components(factors/features) let us plot (X against y)

```
In [29]: #Again :
    X_test
```

```

Out[29]:      last_evaluation  number_project  average_monthly_hours  \
14667          0.93             5             223
13373          0.89             4             137
9246           0.55             5             121
12261          0.97             5             263
9359           0.54             3             271
2419           0.66             6             164
11647          0.63             4             104
13380          0.83             5             216
824            0.56             2             138
14026          0.81             4             179
9008           0.54             3             159
12126          0.49             2             132
811            0.97             7             288
6432           0.67             2             136
5145           0.92             2             198
9546           0.95             4             137
12330          0.57             2             140
3371           0.52             5             222
181            0.84             6             261
9228           0.49             3             267
6520           0.70             4             221
5076           0.71             5             222
6163           0.84             3             239
7981           0.89             4             255
7128           0.99             5             208
2819           0.37             2             159
13084          0.71             4             268
4388           0.96             4             143
9197           0.65             3             183
1291           0.90             6             272
...           ...             ...             ...
9210           0.88             5             223
343            0.56             2             143
2031           0.57             2             160
12730          0.47             2             128
7305           0.67             3             113
2380           0.60             3             205
8912           0.53             4             181
1851           0.53             2             147
9529           0.88             5             225
4124           0.97             3             199
11594          0.55             3             271
12935          0.62             3             152
14351          0.96             6             245
3684           0.56             4             214
14646          0.47             2             135
8283           0.96             4             287

```

12430	0.65	5	195
7600	0.99	4	152
10151	0.73	3	195
4979	0.36	4	97
7516	0.97	4	264
1167	0.50	2	127
11364	0.77	3	144
12793	0.50	4	156
12464	0.98	5	234
8714	0.82	4	190
6409	0.75	4	263
9903	0.87	4	263
4778	0.90	3	142
11310	0.74	5	243

	time_spend_company
14667	5
13373	3
9246	3
12261	5
9359	3
2419	5
11647	7
13380	4
824	3
14026	3
9008	3
12126	3
811	4
6432	6
5145	2
9546	4
12330	3
3371	2
181	4
9228	3
6520	5
5076	3
6163	3
7981	3
7128	2
2819	6
13084	3
4388	3
9197	3
1291	5
...	...
9210	3

343	3
2031	4
12730	3
7305	2
2380	6
8912	3
1851	3
9529	2
4124	3
11594	7
12935	6
14351	4
3684	2
14646	3
8283	5
12430	6
7600	4
10151	2
4979	4
7516	3
1167	3
11364	3
12793	2
12464	5
8714	5
6409	5
9903	2
4778	3
11310	2


[4500 rows x 4 columns]

```
In [30]: X_reduced = pca.transform(X_test)
X_reduced
```

```
Out[30]: array([[ 21.96652508],
                [-64.04607826],
                [-80.03520079],
                ...,
                [ 61.94250711],
                [-59.05667198],
                [ 41.95387206]])
```

```
In [31]: plt.scatter(X_reduced, y_test, color='black')
```


```
Out[31]: <matplotlib.collections.PathCollection at 0x7f7db1094a20>
```



```
In [32]: plt.scatter(X_reduced, y_test, color='black')
plt.plot(X_reduced, test_predicted, color='blue',linewidth=1)
plt.plot(X_reduced, test_predicted, color='red',linewidth=1)

#plt.xticks(())
#plt.yticks(())


plt.show()
```



1.6 Not very insightful? Let us discuss this in class

In [33]: `plt.plot(y_test, test_predicted, 'ro-')`

Out[33]: [`<matplotlib.lines.Line2D at 0x7f7db0ff12b0>`]




```
In [34]: np.std(np.abs(y_test-test_predicted))
```

```
Out[34]: 0.13005465030212102
```


```
In [35]: data4=pd.DataFrame({'actual':y_test,'pred':test_predicted})  
data4.head()
```

```
Out[35]:
```

	actual	pred
14667	0.86	0.601516
13373	0.78	0.649397
9246	0.27	0.523364
12261	0.82	0.615714
9359	0.97	0.613623

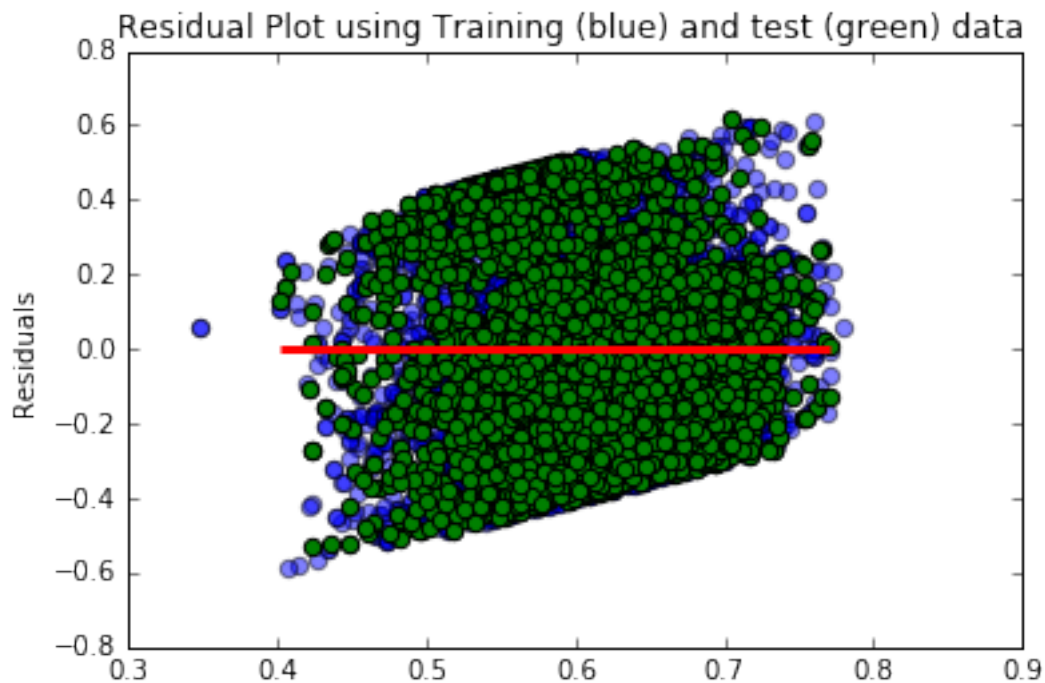
```
In [36]: data4.sort_values('actual').plot(kind='line',x='actual',y='pred')
```

```
Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x7f7db0ff1f98>
```



```
In [37]: plt.scatter(reg.predict(X_train), reg.predict(X_train)-y_train,c='b',s=40,alpha=0.5)  
plt.scatter(reg.predict(X_test),reg.predict(X_test)-y_test,c='g',s=40)  
plt.hlines(y=0,xmin=np.min(reg.predict(X_test)),xmax=np.max(reg.predict(X_test)),color=  
plt.title('Residual Plot using Training (blue) and test (green) data ')  
plt.ylabel('Residuals')
```

Out[37]: <matplotlib.text.Text at 0x7f7db0f04438>



In [38]: data.corr()

Out[38]:

	satisfaction_level	last_evaluation	number_project	\
satisfaction_level	1.000000	0.105021	-0.142970	
last_evaluation	0.105021	1.000000	0.349333	
number_project	-0.142970	0.349333	1.000000	
average_monthly_hours	-0.020048	0.339742	0.417211	
time_spend_company	-0.100866	0.131591	0.196786	
Work_accident	0.058697	-0.007104	-0.004741	
left	-0.388375	0.006567	0.023787	
promotion_last_5years	0.025605	-0.008684	-0.006064	

	average_monthly_hours	time_spend_company	\
satisfaction_level	-0.020048	-0.100866	
last_evaluation	0.339742	0.131591	
number_project	0.417211	0.196786	
average_monthly_hours	1.000000	0.127755	
time_spend_company	0.127755	1.000000	
Work_accident	-0.010143	0.002120	
left	0.071287	0.144822	
promotion_last_5years	-0.003544	0.067433	

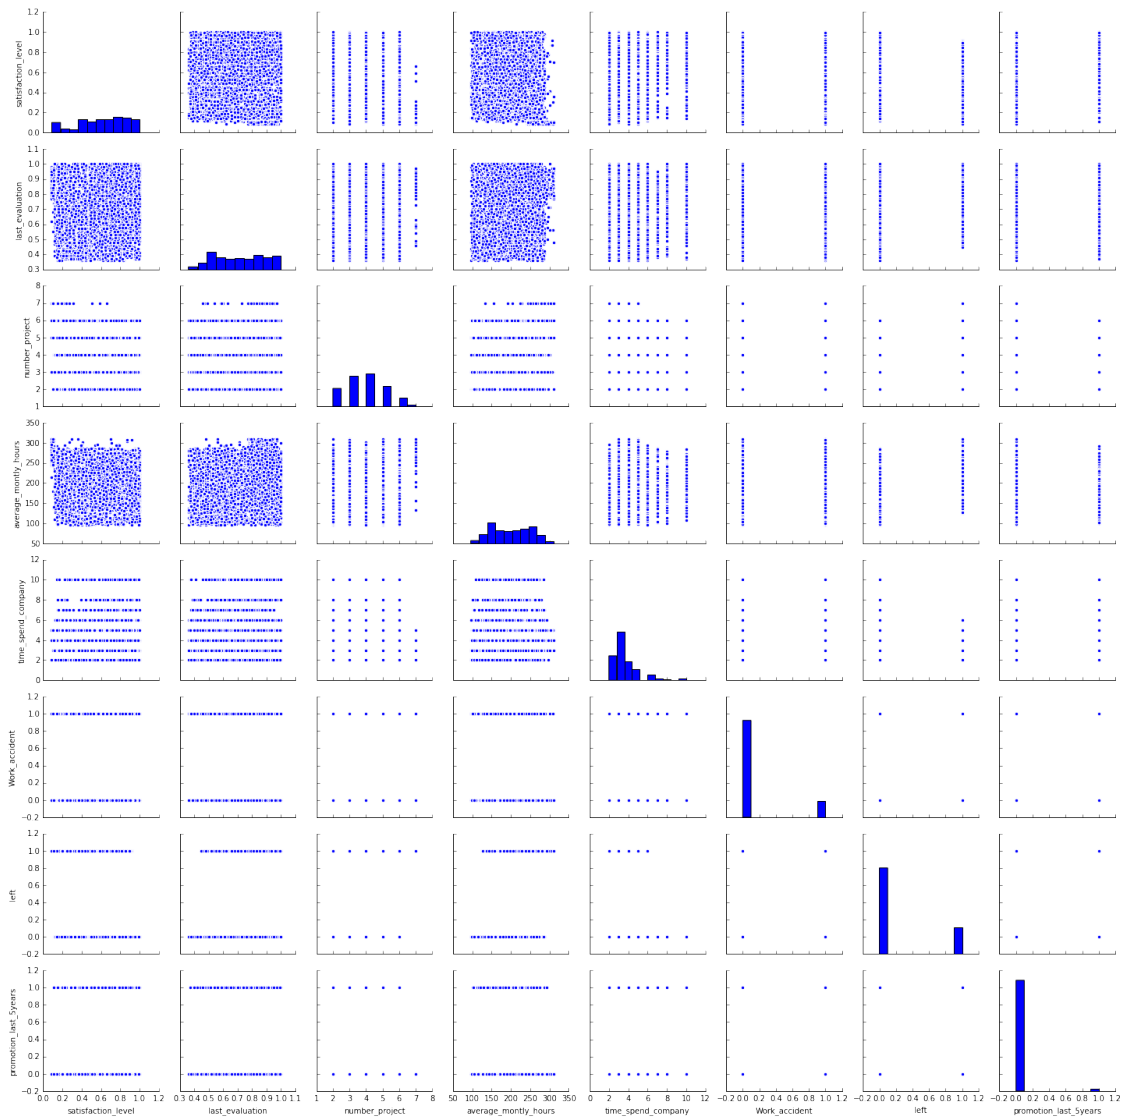
Work_accident	left	promotion_last_5years
---------------	------	-----------------------

satisfaction_level	0.058697	-0.388375	0.025605
last_evaluation	-0.007104	0.006567	-0.008684
number_project	-0.004741	0.023787	-0.006064
average_montly_hours	-0.010143	0.071287	-0.003544
time_spend_company	0.002120	0.144822	0.067433
Work_accident	1.000000	-0.154622	0.039245
left	-0.154622	1.000000	-0.061788
promotion_last_5years	0.039245	-0.061788	1.000000

In [39]: `import seaborn as sns`

In [40]: `sns.pairplot(data)`

Out[40]: `<seaborn.axisgrid.PairGrid at 0x7f7db0ef4128>`



```
In [42]: rng = np.random.RandomState(1)
x = 10 * rng.rand(50)
y = 2 * x - 5 + rng.randn(50)
plt.scatter(x, y);
```

