

Week 4 - Regression Analysis - Python

September 26, 2018

1 Data Warehousing and Data Mining

1.1 Labs

1.1.1 Prepared by Gilroy Gordon

Contact Information SCIT ext. 3643

ggordonutech@gmail.com

gilroy.gordon@utech.edu.jm

1.1.2 Week 3 - Regression Analysis in Python

Additional Reference Resources:

http://scikit-learn.org/stable/modules/linear_model.html

1.2 Objectives

- > Data Transformation
- > Data Mining
 - > Linear Regression
- > Model Evaluation and Prediction
 - > Train/Test Split - 70/30
- > Presentation
 - > Scatter Plot

1.3 Import required libraries and acquire data

NB. The data required was retrieved from the required text for this course. This should assist you in following the concepts from the book better

```
In [2]: # import required libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [3]: data_path = './data/Creditcardprom.csv' # Path to data file
data = pd.read_csv(data_path)
data
```

```
Out[3]:
```

	Income Range	Magazine Promo	Watch Promo	Life Ins Promo	Credit Card Ins.	\
0	40-50,000	Yes	No	No	No	
1	30-40,000	Yes	Yes	Yes	No	
2	40-50,000	No	No	No	No	
3	30-40,000	Yes	Yes	Yes	Yes	
4	50-60,000	Yes	No	Yes	No	
5	20-30,000	No	No	No	No	
6	30-40,000	Yes	No	Yes	Yes	
7	20-30,000	No	Yes	No	No	
8	30-40,000	Yes	No	No	No	
9	30-40,000	Yes	Yes	Yes	No	
10	40-50,000	No	Yes	Yes	No	
11	20-30,000	No	Yes	Yes	No	
12	50-60,000	Yes	Yes	Yes	No	
13	40-50,000	No	Yes	No	No	
14	20-30,000	No	No	Yes	Yes	

	Sex	Age
0	Male	45
1	Female	40
2	Male	42
3	Male	43
4	Female	38
5	Female	55
6	Male	35
7	Male	27
8	Male	43
9	Female	41
10	Female	43
11	Male	29
12	Female	39
13	Male	55
14	Female	19

```
In [4]: # What columns are in the data set ? Do they have spaces that I should consider
data.columns
```

```
Out[4]: Index(['Income Range', 'Magazine Promo', 'Watch Promo', 'Life Ins Promo',
               'Credit Card Ins.', 'Sex', 'Age'],
              dtype='object')
```

1.4 Aim: Can we determine a person's age based on the other factors?

age = a(Income Range) + b(Magazine Promo) + c(Watch Promo) + d(Life Ins Promo)+ e(Credit Card Ins.) + f(Sex)

The coefficients a-f, what are they? What is the relationship between the variables? Does multicollinearity exist?

I have created a function below `create_label_encoder_dict` to assist with this. The function accepts a dataframe object and uses the `LabelEncoder` class from `sklearn.preprocessing` to encode (dummy encoding) or transform non-numerical columns to numbers. Finally it returns a dictionary object of all the encoders created for each column.

The `LabelEncoder` is a useful resource as it not only automatically transforms all values in a column but also keeps a track of what values were transformed from. i.e. It will change all Female to 0 and all Male to 1

```
In [7]: def create_label_encoder_dict(df):
        from sklearn.preprocessing import LabelEncoder

        label_encoder_dict = {}
        for column in df.columns:
            # Only create encoder for categorical data types
            if not np.issubdtype(df[column].dtype, np.number) and column != 'Age':
                label_encoder_dict[column] = LabelEncoder().fit(df[column])
        return label_encoder_dict

In [9]: label_encoders = create_label_encoder_dict(data)
        print("Encoded Values for each Label")
        print("="*32)
        for column in label_encoders:
            print("="*32)
            print('Encoder(%s) = %s' % (column, label_encoders[column].classes_ ))
            print(pd.DataFrame([range(0, len(label_encoders[column].classes_))], columns=label_en
```

Encoded Values for each Label

=====

Encoder(Sex) = ['Female' 'Male']

Encoded Values

Female 0

Male 1

=====

Encoder(Income Range) = ['20-30,000' '30-40,000' '40-50,000' '50-60,000']

Encoded Values

20-30,000 0

30-40,000 1

40-50,000 2

50-60,000 3

=====

Encoder(Magazine Promo) = ['No' 'Yes']

Encoded Values

No 0

Yes 1

=====

```
Encoder(Watch Promo) = ['No' 'Yes']
```

```
    Encoded Values
```

```
No          0
```

```
Yes         1
```

```
=====
```

```
Encoder(Life Ins Promo) = ['No' 'Yes']
```

```
    Encoded Values
```

```
No          0
```

```
Yes         1
```

```
=====
```

```
Encoder(Credit Card Ins.) = ['No' 'Yes']
```

```
    Encoded Values
```

```
No          0
```

```
Yes         1
```

```
In [11]: # Apply each encoder to the data set to obtain transformed values
data2 = data.copy() # create copy of initial data set
for column in data2.columns:
    if column in label_encoders:
        data2[column] = label_encoders[column].transform(data2[column])

print("Transformed data set")
print("="*32)
data2
```

```
Transformed data set
```

```
=====
```

```
Out[11]:
```

	Income Range	Magazine Promo	Watch Promo	Life Ins Promo	\
0	2	1	0	0	
1	1	1	1	1	
2	2	0	0	0	
3	1	1	1	1	
4	3	1	0	1	
5	0	0	0	0	
6	1	1	0	1	
7	0	0	1	0	
8	1	1	0	0	
9	1	1	1	1	
10	2	0	1	1	
11	0	0	1	1	
12	3	1	1	1	
13	2	0	1	0	
14	0	0	0	1	

```
Credit Card Ins.  Sex  Age
```

0	0	1	45
1	0	0	40
2	0	1	42
3	1	1	43
4	0	0	38
5	0	0	55
6	1	1	35
7	0	1	27
8	0	1	43
9	0	0	41
10	0	0	43
11	0	1	29
12	0	0	39
13	0	1	55
14	1	0	19

```
In [16]: # separate our data into dependent (Y) and independent(X) variables
X_data = data2[['Income Range', 'Sex', 'Life Ins Promo', 'Magazine Promo', 'Watch Promo']]
Y_data = data2['Age']
```

1.5 70/30 Train Test Split

We will split the data using a 70/30 split. i.e. 70% of the data will be randomly chosen to train the model and 30% will be used to evaluate the model

```
In [17]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X_data, Y_data, test_size=0.30)
```

```
In [14]: # Import linear model package (has several regression classes)
from sklearn import linear_model
```

```
In [15]: # Create an instance of linear regression
reg = linear_model.LinearRegression()
```

```
In [32]: reg.fit(X_train,y_train)
```

```
Out[32]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

```
In [33]: reg.coef_
```

```
Out[33]: array([ 4.20886076, -8.43037975, -17.51898734,  5.5          ,
                8.26582278])
```

```
In [34]: X_train.columns
```

```
Out[34]: Index(['Income Range', 'Sex', 'Life Ins Promo', 'Magazine Promo',
                'Watch Promo'],
                dtype='object')
```

```
In [35]: print("Regression Coefficients")
pd.DataFrame(reg.coef_,index=X_train.columns,columns=["Coefficient"])
```

Regression Coefficients

```
Out[35]:
```

	Coefficient
Income Range	4.208861
Sex	-8.430380
Life Ins Promo	-17.518987
Magazine Promo	5.500000
Watch Promo	8.265823

```
In [36]: # Intercept
reg.intercept_
```

```
Out[36]: 42.658227848101255
```

```
In [42]: # Make predictions using the testing set
test_predicted = reg.predict(X_test)
test_predicted
```

```
Out[42]: array([26.41772152, 48.14556962, 51.53164557, 43.11392405, 42.64556962])
```

```
In [43]: from sklearn.metrics import mean_squared_error, r2_score
```

```
In [44]: # The mean squared error
print("Mean squared error: %.2f" % mean_squared_error(y_test, test_predicted))
```

```
Mean squared error: 50.14
```

```
In [45]: # Explained variance score: 1 is perfect prediction
# R squared
print('Variance score: %.2f' % r2_score(y_test, test_predicted))
```

```
Variance score: -3.57
```

```
In [47]: help(reg.score)
```

Help on method score in module sklearn.base:

score(X, y, sample_weight=None) method of sklearn.linear_model.base.LinearRegression instance
Returns the coefficient of determination R^2 of the prediction.

The coefficient R^2 is defined as $(1 - u/v)$, where u is the residual sum of squares $((y_{\text{true}} - y_{\text{pred}}) ** 2).sum()$ and v is the total sum of squares $((y_{\text{true}} - y_{\text{true}.mean()}) ** 2).sum()$.

The best possible score is 1.0 and it can be negative (because the model can be arbitrarily worse). A constant model that always predicts the expected value of y , disregarding the input features, would get a R^2 score of 0.0.

Parameters

X : array-like, shape = (n_samples, n_features)
Test samples.

y : array-like, shape = (n_samples) or (n_samples, n_outputs)
True values for X.

sample_weight : array-like, shape = [n_samples], optional
Sample weights.

Returns

score : float
 R^2 of self.predict(X) wrt. y.

```
In [48]: reg.score(X_test,y_test)
```

```
Out[48]: -3.5749185396313994
```

1.5.1 Visualizations

It's difficult to plot a scatter plot with so many dimensions

How about Dimensionality Reduction?

One such method - Principal Component Analysis

```
In [50]: from sklearn.decomposition import PCA
```

```
In [51]: pca = PCA(n_components=1)
```

```
In [53]: pca.fit(data2[X_train.columns])
```

```
Out[53]: PCA(copy=True, iterated_power='auto', n_components=1, random_state=None,
          svd_solver='auto', tol=0.0, whiten=False)
```

```
In [54]: pca.components_
```

```
Out[54]: array([[ 0.96117326, -0.10803145,  0.08602344,  0.23885282, -0.0049457 ]])
```

```
In [55]: pca.n_features_
```

```
Out[55]: 5
```

```
In [56]: pca.n_components_
```

```
Out[56]: 1
```

Now that we can reduce our components(factors/features) let us plot (X against y)

```
In [59]: #Again :  
X_test
```

```
Out[59]:
```


	Income	Range	Sex	Life	Ins	Promo	Magazine	Promo	Watch	Promo
6		1	1			1		1		0
0		2	1			0		1		0
12		3	0			1		1		1
1		1	0			1		1		1
2		2	1			0		0		0

```
In [60]: X_reduced = pca.transform(X_test)  
X_reduced
```

```
Out[60]: array([[ -0.15821581],  
                [ 0.71693401],  
                [ 1.86721646],  
                [-0.05513007],  
                [ 0.47808119]])
```

```
In [61]: plt.scatter(X_reduced, y_test, color='black')
```


```
Out[61]: <matplotlib.collections.PathCollection at 0x7f402752ed68>
```




```
In [64]: plt.scatter(X_reduced, y_test, color='black')
plt.scatter(X_reduced, test_predicted, color='blue')
plt.plot(X_reduced, test_predicted, color='red',linewidth=1)

plt.xticks(())
plt.yticks(())

plt.show()
```



1.6 Not very insightful? Let us discuss this in class